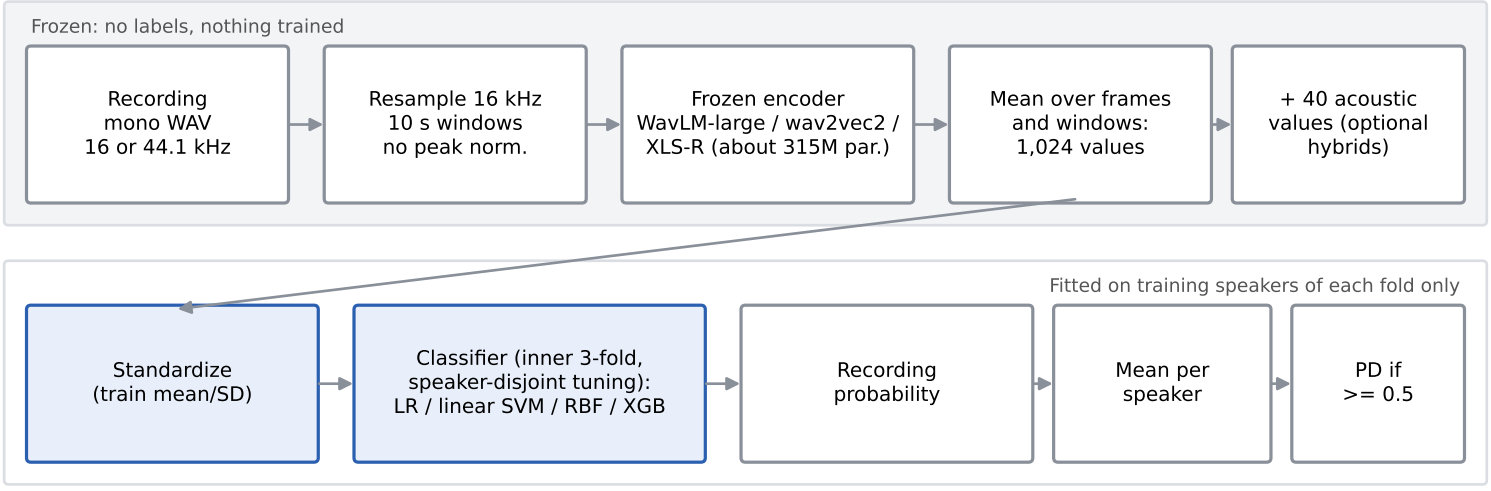
Frozen: no labels, nothing trained
Recording
mono WAV
16 or 44.1 kHz
Resample 16 kHz
10 s windows
no peak norm.
Frozen encoder
WavLM-large / wav2vec2 /
XLS-R (about 315M par.)
Mean over frames
and windows:
1,024 values
+ 40 acoustic
values (optional
hybrids)
Fitted on training speakers of each fold only
Standardize
(train mean/SD)
Classifier (inner 3-fold,
speaker-disjoint tuning):
LR / linear SVM / RBF / XGB
Recording
probability
Mean per
speaker
PD if
>= 0.5