Outer loop: 47 speakers in 5 folds (one fold = 9 or 10 people, both classes)

repeat 5 times, each fold is the test fold once

Inner loop (inside the 4 training folds only): 3 speaker-disjoint inner folds

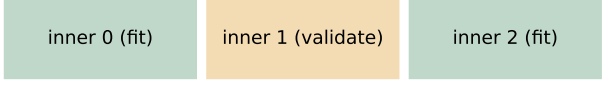

1. For every configuration of the grid: fit on 2 inner folds, score the 3rd; rotate.
2. Pool the inner predictions per speaker; compute speaker-level AUPRC; keep the best configuration.
3. Refit that configuration on all training speakers (and redo feature selection if used).
4. Predict the untouched test fold once; average each test speaker's recordings; threshold 0.5.

Nothing learned (scaling, feature selection, hyperparameters, model) ever sees a test speaker.